

identity after assembly; each such control section can be loaded and relocated independently of the others. Different control sections are most often used to subdivide a program into subroutines or other logical subdivisions of a program. The programmer can assemble, load, and manipulate each of these control sections separately. The resulting flexibility is a major benefit of using control sections. We consider examples of this when we discuss linkage editors in Chapter 3.

When control sections form logically related parts of a program, it is necessary to provide some means for *linking* them together. For example, instructions in one control section might need to refer to instructions or data located in another section. Because control sections are independently loaded and relocated, the assembler is unable to process these references in the usual way. The assembler has no idea where any other control section will be located at execution time. Such references between control sections are called *external references*. The assembler generates information for each external reference that will allow the loader to perform the required linking. In this section we describe how external references are handled by our assembler. Chapter 3 discusses in detail how the actual linking is performed.

Figure 2.15 shows our example program as it might be written using multiple control sections. In this case there are three control sections: one for the main program and one for each subroutine. The START statement identifies the beginning of the assembly and gives a name (COPY) to the first control section. The first section continues until the CSECT statement on line 109. This assembler directive signals the start of a new control section named RDREC. Similarly, the CSECT statement on line 193 begins the control section named WRREC. The assembler establishes a separate location counter (beginning at 0) for each control section, just as it does for program blocks.

Control sections differ from program blocks in that they are handled separately by the assembler. (It is not even necessary for all control sections in a program to be assembled at the same time.) Symbols that are defined in one control section may not be used directly by another control section; they must be identified as external references for the loader to handle. Figure 2.15 shows the use of two assembler directives to identify such references: EXTDEF (external definition) and EXTREF (external reference). The EXTDEF statement in a control section names symbols, called *external symbols*, that are defined in this control section and may be used by other sections. Control section names (in this case COPY, RDREC, and WRREC) do not need to be named in an EXTDEF statement because they are automatically considered to be external symbols. The EXTREF statement names symbols that are used in this control section and are defined elsewhere. For example, the symbols BUFFER, BUFEND, and LENGTH are defined in the control section named COPY and made available to the other sections by the EXTDEF statement on line 6. The third control section (WRREC) uses two of these symbols, as specified in its EXTREF statement.

Source statement			
6	START	0	COPY FILE FROM INPUT TO OUTPUT
	EXTDEF	BUFFER, BUFEND, LENGTH	
	EXTREF	RDREC, WRREC	
10	RETADR	RETADR	SAVE RETURN ADDRESS
11	RTL	RDREC	READ INPUT RECORD
12	+JBUB	LENGTH	TEST FOR EOF (LENGTH = 0)
13	LDA	#0	
14	COMP	EMDFIL	EXIT IF EOF FOUND
15	JEQ	WRREC	WRITE OUTPUT RECORD
16	+JSUB	CLOOP	LOOP
17	J	=C'EOF'	INSERT END OF FILE MARKER
18	LDA	BUFFER	
19	STA	#3	SET LENGTH = 3
20	LDA	LENGTH	
21	STA	WRREC	
22	+JBUB	J @RETADR	RETURN TO CALLER
23	RETADR	RESW	LENGTH OF RECORD
24	LENGTH	1	
25	RESW	1	
26	LTORQ		
27	BUFFER	4096	4096-BYTE BUFFER AREA
28	BUFEND	*	
29	MAXLEN	EQU	BUFEND-BUFFER
30	RDREC	EQU	
31	CSECT		
32			SUBROUTINE TO READ RECORD INTO BUFFER
33	EXTREF	BUFFER, LENGTH, BUFEND	
34	CLEAR	X	CLEAR LOOP COUNTER
35	CLEAR	A	CLEAR A TO ZERO
36	CLEAR	S	CLEAR S TO ZERO
37	LDT	MAXLEN	
38	TD	INPUT	TEST INPUT DEVICE
39	JEQ	RLOOP	LOOP UNTIL READY
40	RD	INPUT	READ CHARACTER INTO REGISTERS
41	COMPR	A, S	TEST FOR END OF RECORD (X)
42	JEQ	EXIT	EXIT LOOP IF EOR
43	+STCH	BUFFER, X	STORE CHARACTER IN BUFFER
44	TIXR	T	LOOP UNLESS MAX LENGTH
45	JLT	RLOOP	HAS BEEN REACHED
46	+STX	LENGTH	SAVE RECORD LENGTH
47	RSUB		RETURN TO CALLER
48	INPUT	X'F1'	CODE FOR INPUT DEVICE
49	MAXLEN	WORD	BUFEND-BUFFER
50	WRREC	CSECT	
51			SUBROUTINE TO WRITE RECORD FROM BUFFER
52	EXTREF	LENGTH, BUFFER	
53	CLEAR	X	CLEAR LOOP COUNTER
54	+LDT	LENGTH	
55	TD	=X'05'	TEST OUTPUT DEVICE
56	JEQ	WLOOP	LOOP UNTIL READY
57	+LDCH	BUFFER, X	GET CHARACTER FROM BUFFER
58	WD	=X'05'	WRITE CHARACTER
59	TIXR	T	LOOP UNTIL ALL CHARACTERS
60	JLT	WLOOP	HAVE BEEN WRITTEN
61	RSUB		RETURN TO CALLER
62	END	FIRST	

Figure 2.15 Illustration of control sections and program linking.

Now we are ready to look at how external references are handled by the assembler. Figure 2.16 shows the generated object code for each statement in the program. Consider first the instruction

15 0003 CLOOP +JSUB RDREC 4B100000

The operand (RDREC) is named in the EXTREF statement for the control section, so this is an external reference. The assembler has no idea where the control section containing RDREC will be loaded, so it cannot assemble the address for this instruction. Instead the assembler inserts an address of zero and passes information to the loader, which will cause the proper address to be inserted at load time. The address of RDREC will have no predictable relationship to anything in this control section; therefore relative addressing is not possible. Thus an extended format instruction must be used to provide room for the actual address to be inserted. This is true of any instruction whose operand involves an external reference.

Similarly, the instruction

160 0017 +STCH BUFFER, X 57900000

makes an external reference to BUFFER. The instruction is assembled using extended format with an address of zero. The *x* bit is set to 1 to indicate indexed addressing, as specified by the instruction. The statement

190 0028 MAXLEN WORD BUFEND-BUFFER 000000

is only slightly different. Here the value of the data word to be generated is specified by an expression involving two external references: BUFEND and BUFFER. As before, the assembler stores this value as zero. When the program is loaded, the loader will add to this data area the address of BUFEND and subtract from it the address of BUFFER, which results in the desired value.

Note the difference between the handling of the

Note the difference between the handling of the expression on line 190 and the similar expression on line 107. The symbols BUFEND and BUFFER are defined in the same control section with the EQU statement on line 107. Thus the value of the expression can be calculated immediately by the assembler. This could not be done for line 190; BUFEND and BUFFER are defined in another control section, so their values are unknown at assembly time.

As we can see from the above discussion, the assembler must remember (via entries in SYMTAB) in which control section a symbol is defined. Any attempt to refer to a symbol in another control section must be flagged as an error unless the symbol is identified (using EXTREF) as an external reference. The assembler must also allow the same symbol to be used in different control

Source statement Objec

	0000	START	0
		EXTREF	BUFFER, BUFEND, LENGTH
		RETADR	RDRREC, WRREC
		STL	RETADR
		LJUB	RDRREC
		LDA	LENGTH
		COMP	#0
		JEQ	ENDFIL
		+JSUB	WRREC
		J	CLOOP
		LDA	=C'EOF'
		STA	BUFFER
		LDA	#3
		STA	LENGTH
		+JSUB	WRREC
		J	GRETADR
		RETADR	RESW
		LENGTH	RESW
			LTORG
		*	=C'EOF'
	0003	BUFFER	RESB
		BUFEND	4096
	1000	MAXLEN	EQU
			BUFEND-BUFFER
0000	RDREC	CSECT	
SUBROUTINE TO READ RECORD INTO B			
		EXTREF	BUFFER, LENGTH, BUFEND
		CLEAR	X
		CLEAR	A
		CLEAR	S
		LDT	MAXLEN
		TD	INPUT
		RLOOP	TD
	000C	JEQ	RLOOP
	000F	RD	INPUT
	0012	COMPR	A,S
	0014	JEQ	EXIT
	0017	+STCH	BUFFER,X
	001B	TIXR	T
	001D	JLT	RLOOP
	0020	EXIT	+STX
	0024	RSUB	LENGTH
	0027	INPUT	BYTE
	0028	MAXLEN	WORD
			X'F1'
			BUFEND-BUFFER
0000	WRREC	CSECT	
SUBROUTINE TO WRITE RECORD FROM E			
		EXTREF	LENGTH, BUFFER
		CLEAR	X
		+LDT	LENGTH
	0002	TD	=X'05'
	0006	WLOOP	WLOOP
	0009	JEQ	+LDCH
	000C	WD	BUFFER,X
	0010	TIXR	=X'05'
	0013	JLT	WD
	0015	RSUB	T
	0018	END	WLOOP
	001B	*	FIRST
			=X'05'

Figure 2.16 Program from Fig. 2.15 with object code

sections. For example, the conflicting definitions of MAXLEN on lines 107 and 190 should cause no problem. A reference to MAXLEN in the control section COPY would use the definition on line 107, whereas a reference to MAXLEN in RDREC would use the definition on line 190.

So far we have seen how the assembler leaves room in the object code for the values of external symbols. The assembler must also include information in the object program that will cause the loader to insert the proper values where they are required. We need two new record types in the object program and a change in a previously defined record type. As before, the exact format of these records is arbitrary; however, the same information must be passed to the loader in some form.

The two new record types are Define and Refer. A Define record gives information about external symbols that are defined in this control section—that is, symbols named by EXTDEF. A Refer record lists symbols that are used as external references by the control section—that is, symbols named by EXTREF. The formats of these records are as follows.

Define record:

Col. 1	D*
Col. 2-7	Name of external symbol defined in this control section
Col. 8-13	Relative address of symbol within this control section (hexadecimal)
Col. 14-73	Repeat information in Col. 2-13 for other external symbols

Refer record:

Col. 1	R
Col. 2-7	Name of external symbol referred to in this control section
Col. 8-23	Names of other external reference symbols

The other information needed for program linking is added to the Modification record type. The new format is as follows.

Modification record (revised):

Col. 1	M
Col. 2-7	Starting address of the field to be modified, relative to the beginning of the control section (hexadecimal)
Col. 8-9	Length of the field to be modified, in half-bytes (hexadecimal)

Col. 10 Modification flag (+ or -)
 Col. 11-16 External symbol whose value is to be added to or subtracted from the indicated field

The first three items in this record are the same as previously discussed. The new items specify the modification to be performed: adding or subtracting the value of some external symbol. The symbol used for modification is defined either in this control section or in another one.

Figure 2.17 shows the object program corresponding to the source code in Figure 2.16. Notice that there is a separate set of object program records.

00000001033
BUFEND001033 LENGTH00002D
WRREC
1D1720274B10000032023900003320074B100003F2FEC032016
1D0D0100030F200A4B1000003E2000
3003454F46
405+RDREC
1105+WRREC
2405+WRREC
0000

Figure 2-17 Object program corresponding to Fig. 2-15.

Header through End) for each control section. The records for each control section are exactly the same as they would be if the sections were assembled separately.

The Define and Refer records for each control section include the symbols named in the EXTDEF and EXTREF statements. In the case of Define, the record also indicates the relative address of each external symbol within the control section. For EXTREF symbols, no address information is available. These symbols are simply named in the Refer record.

Now let us examine the process involved in linking up external references, beginning with the source statements we discussed previously. The address field for the JSUB instruction on line 15 begins at relative address 0004. Its initial value in the object program is zero. The Modification record

M00000405+RDREC

in control section COPY specifies that the address of RDREC is to be added to this field, thus producing the correct machine instruction for execution. The other two Modification records in COPY perform similar functions for the instructions on lines 35 and 65. Likewise, the first Modification record in control section RDREC fills in the proper address for the external reference on line 160.

The handling of the data word generated by line 190 is only slightly different. The value of this word is to be BUFEND-BUFFER, where both BUFEND and BUFFER are defined in another control section. The assembler generates an initial value of zero for this word (located at relative address 0028 within control section RDREC). The last two Modification records in RDREC direct that the address of BUFEND be added to this field, and the address of BUFFER be subtracted from it. This computation, performed at load time, results in the desired value for the data word.

In Chapter 3 we discuss in detail how the required modifications are performed by the loader. At this time, however, you should be sure that you understand the concepts involved in the linking process. You should carefully examine the other Modification records in Fig. 2.17, and reconstruct for yourself how they were generated from the source program statements.

Note that the revised Modification record may still be used to perform program relocation. In the case of relocation, the modification required is adding the beginning address of the control section to certain fields in the object program. The symbol used as the name of the control section has as its value the required address. Since the control section name is automatically an external symbol, it is available for use in Modification records. Thus, for example, the Modification records from Fig. 2.8 are changed from

M00000705
M00001405
M00002705

M00000705+COPY
M00001405+COPY
M00002705+COPY

In this way, exactly the same mechanism can be used for program re-linking and for program linking. There are more examples in the next chapter.

The existence of multiple control sections that can be relocated independently of one another makes the handling of expressions slightly more complicated. Our earlier definitions required that all of the relative terms in an expression be paired (for an absolute expression), or that all except one term be paired (for a relative expression). We must now extend this restriction to require that both terms in each pair must be relative within the same control section. The reason is simple—if the two terms represent relative locations in different control sections, their difference is an absolute value (regardless of the control section in which the expression is located). On the other hand, if they are in different control sections, their difference has a value that is unpredictable (and therefore probably useless). For example, the expression

BUFEND-BUFFER

has as its value the length of BUFFER in bytes. On the other hand, the expression

RDREC-COPY

has as its value the difference in the load addresses of the two control sections. This depends on the way run-time storage is allocated; it is unlikely to be predictable to anyone other than the application program.

When an expression involves external references, the assembler cannot generally determine whether or not the expression is legal. The pairing of terms to test legality cannot be done without knowing which of the terms are in the same control sections, and this is unknown at assembly time. Instead, the assembler evaluates all of the terms it can, and combines them to an initial expression value. It also generates Modification records so that it can finish the evaluation. The loader can then check the expression for legality. We discuss this further in Chapter 3 when we examine the design of a linker.